

ConfiBERT: A Language Model for Political Conflict*

Patrick T. Brandt; UT Dallas, pbrandt@utdallas.edu

Sultan Alsarra; King Saud University, salsarra@ksu.edu.sa

Vito J. D’Orazio; West Virginia University, vito.dorazio@mail.wvu.edu

Dagmar Heintze; UT Dallas, dagmar.heintze@utdallas.edu

Latifur Khan; UT Dallas, lkhan@utdallas.edu

Shreyas Meher; UT Dallas, shreyas.meher@utdallas.edu

Javier Osorio; University of Arizona, josorio1@arizona.edu

Marcus Sianan; UT Dallas, marcus.sianan@utdallas.edu

December 20, 2024

Abstract

Conflict scholars have used rule-based approaches to extract information about political violence from news reports and texts. Recent Natural Language Processing developments move beyond rigid rule-based approaches. We review our recent ConfiBERT language model (Hu et al. 2022) to process political and violence related texts. The model can be used to extract actor and action classifications from texts about political conflict. When fine-tuned, results show that ConfiBERT has superior performance in accuracy, precision and recall over other large language models (LLM) like Google’s Gemma 2 (9B), Meta’s Llama 3.1 (7B), and Alibaba’s Qwen 2.5 (14B) within its relevant domains. It is also hundreds of times faster than these more generalist LLMs. These results are illustrated using texts from the BBC, re3d, and the Global Terrorism Dataset (GTD).

*Previous versions have been presented at the virtual Methods in Event Detection Colloquium (September 2024). This research was supported by NSF award 2311142; used Delta at NCSA / University of Illinois through allocation CIS220162 from the Advanced Cyberinfrastructure Coordination Ecosystem: Services & Support (ACCESS) program, which is supported by NSF grants 2138259, 2138286, 2138307, 2137603, and 2138296. This material is based upon High Performance Computing (HPC) resources supported by the University of Arizona TRIF, UITS, and Research, Innovation, and Impact (RII). Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation, NCSA, or the affiliated university resource providers. Thanks for suggestions go to Scott Althaus, R. Michael Alvarez, Ben Bagozzi, Rebecca Cordell, Ryan Kennedy, Hyein Kim, Shahryar Minhas, Philip Schrod, and Nora Webb Williams.

1 Introduction

A common text as data problem in the social sciences is the need to migrate information in a large corpora into a data summary of that text. As social (political) scientists we possess replicable and encoded domain expertise to understand these texts. There are components of the *processing* and *summarization* of large corpora that are better understood by experts in machine learning and natural language processing (NLP). How then should one combine the insights of domain experts and computational and information scientists of natural language? Here we present one possible answer using recent large language models to illustrate how domain expertise and computational science can be combined.

The conflict event data we care about are typically from human-coded news reports. The transformation of unstructured news texts into structured "who-did-what-to-whom" event data is fundamental in the international relations, and studies of conflict and political violence. The process of gathering and preparing data for analysis in this domain is: assemble a corpus, filter for relevant information, identify target events, and annotate event attributes. This process can be costly and time consuming: time required for data collection, filtering the large amount of unstructured text, training human annotators on the usage of the applied ontology, and several rounds of quality controls to ultimately achieve a curated text corpus. This parallels the widely systematized way to approach text-as-data in much of international relations and the social sciences (O'Connor et al., 2013; Grimmer et al., 2022). Ultimately, these pre-processing steps are applied to the unstructured corpus to permit the extraction of relevant information on actors (who), actions (did what) and the targets (to whom) for political science research.

This transformation of unstructured news text faces several well-known challenges that researchers in the field have long faced.¹ These can be categorized into three main areas:

1. Filtering for relevant information or reports
2. Identifying events
3. Annotating event attributes

While these areas can be challenging, they align with common tasks in language learning and computational information science. In NLP these latter tasks are addressed through techniques such as semantic role learning (SRL), automated content extraction (ACE). Here we address how these problems can be addressed with a language model called ConfiBERT or a Bidirectional Encoder Representations from Transformers (BERT) model trained on conflict data. ConfiBERT, as a political "conflict" informed model, does this

¹We presume the existence of a relevant corpus of texts in digital form, akin to economists' can-opener assumption.

using a BERT model.² ConflBERT is a language model trained on domain-specific conflict and political violence data and achieves superior performance in the summarization and classification tasks political scientists perform on unstructured text corpora. ConflBERT (for English) is publicly available at Hugging Face, and is trained on an expert-curated corpus of 33.7 GB about conflict and political violence. Combining domain-specific knowledge and computational tools in ConflBERT improves unstructured conflict texts processing. The method discussed here is a domain-specific pretrained language model implemented for English (Hu et al., 2022), and there are variants for Spanish (Yang et al., 2023), and Arabic (Alsarra et al., 2023). ConflBERT, as a political science trained model, can help encode information about events from text without having to have a fully specified ontology of actors or their interactions (a major impediment in earlier domain-based methods).

Our proposed method uses a transformer language models (Devlin et al., 2018) and large amounts of politically relevant news texts as training data. We re-cast the problems of the political science domain into those more commonly seen in the information and computer science domains of text and natural language processing and inferences. This trades human annotation and classification costs for computational resources (which grow more powerful and cheaper), while augmenting the speed of data generation. However, how this is done in the computational linguistics and information sciences community needs to be bridged to how social scientists think about these problems. There the focus is on annotation of spans of text that correspond to linguistic and contextual entities (by human or machine). In contrast, political and social scientists focus on event attributes, their modality and characteristics (Olsen et al., 2024). Similar domain-specific BERT models have shown to equally outperform generic BERT models in text classifications in other scientific fields such as: biomedical SCIBERT (Beltagy et al., 2019), material sciences, MatSCIBERT (Gupta et al., 2022), legal LegalBERT Chalkidis et al. (2020), finance (FinBERT) Araci (2019), clinical notes (ClinicalBERT) (Huang et al., 2020), and patent texts (patentBERT) (Lee and Hsiang, 2019).

For the political science domain we want to use ConflBERT to accomplish three key information extraction and summarization tasks that are part of "coding event data":

Filtering for Relevant Information: In political science, the sheer volume of available text data — from news articles and social media posts to diplomatic cables and policy documents — necessitates efficient filtering. The challenge lies not just in identifying texts that contain political content, but in discerning which texts contain information relevant to specific research questions or event types of interest.

For instance, a researcher must differentiate between reports of actual violent events and discussions

²The conflict domain part is first and the LLM / AI tool is second. Since we started this in 2021, the LLM tool of the time was BERT. However, we further introduce a conflict-specific adaptation of LLAMA (Dubey et al., 2024), a generative AI-based model, which we termed ConflLlama and trained on conflict-specific training data in section 4 to permit a performance comparison between ConflBERT and current generative AI models.

of potential conflicts or historical references. This task becomes even more complex when dealing with subtle forms of political interaction or when trying to capture early warning signs of emerging conflicts. ConflBERT’s binary classification for political violence addresses this challenge by providing a confidence score for the relevance of a given text to conflict or political violence. This allows researchers to quickly sift through large volumes of text, focusing their attention on the most pertinent sources.

Identifying Events: Once relevant texts are identified, the next challenge is to accurately identify and delineate specific events within these texts. This task is complicated by several factors. Multiple events may be described within a single text, or events may be described across multiple sentences or even multiple documents. The language used to describe events can be ambiguous or metaphorical, and different sources may provide conflicting accounts of the same event.

In political science, accurate event identification is crucial for creating reliable event datasets, which form the backbone of many quantitative analyses in the field. Misidentification of events can lead to skewed results and flawed policy recommendations. ConflBERT’s multiclass classification contributes to this task by categorizing the type of event described in the text (e.g., Armed Assault, Bombing or Explosion). This classification can serve as a starting point for more detailed event extraction and coding. But this step often needs iterations and revisions thus requiring speed and computational efficiency as well as accuracy.

Annotating Event Attributes: Perhaps the most challenging aspect in unstructured text processing is the detailed annotation of event attributes—the "who," "what," "to whom," "where," and "when" of each identified event. This requires not just entity recognition, but also understanding the roles these entities play in the event and the relationships between them.

Political scientists often need to code additional attributes beyond the basic event structure. These might include the intensity or scale of the event, the motivations or goals of the actors involved, the immediate and long-term consequences of the event, and the broader political context in which the event occurred. ConflBERT’s Named Entity Recognition (NER) capabilities provide a foundation for this task by identifying key actors, organizations, and locations within the text. However, fully automating the annotation of complex event attributes remains a significant challenge and often requires a combination of machine learning techniques and expert human coding.

We compare our model to some of the latest generative large language models and show its highly favorable performance on the tasks defined above. ConflBERT significantly alleviates the human annotation bottleneck in these three areas. We also review the use of ConflBERT and our related innovations for engineering data from text. ConflBERT was trained on conflict and political violence related data in 2021, and consistently outperforms original BERT (Hu et al., 2022), Electra and RoBERTa (Osorio et al., 2024). Its non-English variants, ConflBERT-Spanish outperforms BERT variants like mBERT and BETO Yang

et al. (2023); and, ConflBERT-Arabic does the same relative to AraBERT (Osorio et al., 2024).

Furthermore, *ConflBERT outperforms current generative AI models such as Meta’s Llama 3.1 (Dubey et al., 2024), and Gemma 2 (Team Gemma et al., 2024) in binary classification and NER tasks.* Here we compare ConflBERT to nascent generative LLMs like Google’s Gemma 2, Meta’s Llama, and Alibaba’s Qwen 2.5 (Hui et al., 2024). We demonstrate that the fine-tuned models used can outperform the more general models in terms of accuracy, precision, recall, and computational time.

ConflBERT has several advantages over comparable contemporaneous methods for machine coding events. First, it is easily deployed and replicable as a method since it is open source and can be deployed on conventional hardware. Second, it is significantly better on comparable, relevant quality metrics and faster than rival or even newer generative artificial intelligence (AI) methods that used encoder-decoder technologies with graphical processing units (GPUs). Third, it can be rapidly deployed to detect new event data and their characteristics.³ This means it can be tuned and adjusted as needed for new cases, data, and texts. This allows users to improve the extraction, coverage (geographically, and as we show, linguistically), across new data and training domains. Fourth this means that additional downstream tasks such as recoding texts, extracting additional variables or features, etc. are all much faster and easier than what has historically been the case. We show this in our examples where differences are seen in the classifications of terrorist event types in the GTD dataset across the LLMs. ConflBERT and domain-specific models provide much better results compared to generalist LLMs like Gemma, Llama, and Qwen.

2 The Event Coding Problem

In event data research, scholars break down texts into key attributes: actors (sources and targets), actions, locations, and dates.⁴ With actor coding, there are two broad approaches that eschew human coding: mining past data to propose new groups or categories of actors (Solaimani et al., 2017b) and machine learning or transformer approaches utilizing BERT-based models (Dai et al., 2022; Hu et al., 2022; Parolin et al., 2022; Alsarra et al., 2023; Yang et al., 2023; Halterman et al., 2023b). To code actions prior work employed sparse parsing using human-annotated dictionaries (Schrodt, 2001; Osorio et al., 2020), whereas newer approaches handle new ontology or action extensions through either up-sampling (Halterman and Radford, 2021), natural language inference (NLI) (Dai et al., 2022; Hu et al., 2022; Lefebvre and Stoeck, 2022; Parolin et al., 2022; Halterman et al., 2023a; Croicu, 2024), or zero-shot prompts (Hu et al., 2024). Geographic coding in earlier

³We thank a reviewer / commentators for emphasis on this point.

⁴For those interested in getting event data and not exploring the weeds of this process, see Halterman et al. (2023a)’s Political Language Ontology for Verifiable Event Records (PLOVER) and the POLitical Event Classification, Attributes, and Types (POLECAT) dataset. which are a record of domestic and international political interactions described in international news reports from 2010 to the present. The news reports are in English or machine translated into English from Arabic, Chinese, French, Portuguese, Russian, or Spanish before they are coded.

work relied on the location inferred from the actors to identify where the event occurred. Some approaches to determine location use sparse parsing Osorio et al. (2020), word embedding and NER (Halterman, 2017; Imani et al., 2017, 2019, e.g.), and even BERT (Halterman et al., 2023a). For date or time coding of events, researchers generally parse the byline of the news report to acquire the publication date (Osorio et al., 2020), but the publication and the event occurrence dates are not always the same. A recent approach is to apply BERT technology to reconcile date information from the news story (Halterman et al., 2023a). All of these information extraction approaches are prone to various errors (Brandt and Sianan, 2024), and latest methods attempt to reduce them using BERT-alike language models.

ConfiBERT provides a domain level solution to a common problem. For most generic generative and extractive artificial intelligence (AI) tasks, an LLM needs broad training. These training steps generate huge costs in terms of 1) training data, 2) human / expert time, and 3) computational complexity to combine and produce the relevant model. In a domain-specific application like ConfiBERT, several choices make these challenges much more feasible for a social science tool. First, creating an extractive LLM or a BERT LLM (or even for that matter a *simple* predictive or generative suggestion model) can be done much more rapidly and cheaply. Since there is domain knowledge and insight provided in the initial training steps, steps 1 and 2 above for training a generic LLM are greatly scaled back, resulting in a superior model in a shorter period of time. Second, ConfiBERT can then be augmented or expanded in various ways (which we demonstrate below) to focus on harder tasks others have attempted, such as ontology extension (Radford, 2021), actor detection and recognition (Solaimani et al., 2017a,b), and image processing applications (Steinert-Threlkeld, 2019; Wen et al., 2021).

The earlier needs to extract actors, action events (verbs) and additional information from texts for political science and international relations studies of conflict are accommodated in three different NLP tasks. The three main tasks that ConfiBERT addresses are:

Classification Which texts contain relevant information about politics, conflict, violence? We give examples of this below based on data from the BBC and re3d text corpora. These are

1. binary classifications: yes / no questions
2. multi-label classifications: in a series of reports about protests, which types of protest are present (labor, peaceful, violent, etc.)?

Named Entity Recognition (NER) What are the “who” and “whom” that characterize the event? These are most typically the linguistic subjects and objects of the sentences and clauses (subject to textual disambiguation and co-referencing). But making sense of them becomes a task for a political scientist

to identify the source / initiator of a political event toward a target or other political actor. We give an example below of doing this for texts about terrorist attacks from GTD.

Masking / Coding new entities and / or events is the extension of any ontology of new kinds of events. This can include teaching a model which events are new ones, ones to be excluded, or newly emergent actors and their roles.

The first two of these tasks are *extractive* tasks often best handled using statistical or machine learning algorithms to make suggestions for validation. These tasks are predictive in the sense that a model is trained to learn and predict patterns based on repeated past examples or interactions. Prior SVM-like versions of similar classification problems for international conflict events have been addressed with simpler models that are context free (D’Orazio et al., 2014). But they are the context-dependent tasks for which BERT-alike models are optimized to do. A method like ConfiBERT improves these previous approaches using longer embedded patterns of related text and their rapid abilities to find additional patterns faster than a human. It is able to accomplish this in situations where the events or entities to be classified are rare (there may be few examples to mask or learn from) or where there is a class imbalance (i.e., the presence of mass atrocities is rare, even among a collection of stories about multiple crimes and political attacks).

The last task is harder, but can begin with an LLM or a BERT-like model. Determining whether another event is like a prior one (or a related class or actor) requires providing examples that omit the thing to be predicted (masking or hiding it) and then seeing how well the model does. This problem can apply to new actors and events and the determination of whether the model / coder is correct relies on the domain knowledge of the social scientist. Further, the identification of new actors is often a masking task for which BERT-alike models are designed.⁵

More generally these are problems of open information extraction (OIE), a critical NLP task. OIE extracts structured information from unstructured source text by detecting relationships without relying on a previously defined ontology. These are generic text and information extraction tasks that are necessary in any coding of texts as data or processing human produced texts. For a political scientist there are *domain specific versions* of these tasks. In the actor identification domain, OIE is a necessary task for accurate event detection and action feature classification.

The problem of finding political actors is twofold. One challenge lies in the ability to automatically extract new political entities from text. The other is to correctly associate these actors with their respective political environment (role, country, organization, and group, etc.). But note this is *not* specific to the domain of political science (remove the word “political” from the first sentences of this paragraph). To place an event

⁵Our recent work in actor detection proposes Distant Supervised (DS) and Zero-Shot (ZS) approaches for extracting political actors and their roles using pre-trained language models (Parolin et al., 2021; Hu et al., 2024).

in context, it is often critical to know not only an actor and their role, but also the actor’s relation to the government (e.g., a rebel group) or other organizations or individuals. Across conflict research, actors are a fundamental unit of analysis: individuals, groups, organizations, countries, etc. Actors have attributes, such as a role within government or an affiliation with a particular group, and the attributes of actors can change over time. Accurately extracting conflict events from text, or building networks of political interactions, depends on our ability to identify actors and their roles. For example, a meeting between Joe Biden and Emmanuel Macron, two heads of state, is categorically different from a meeting between two elected officials from the same country. Actor identification is a domain knowledge encoding of information into labels that are relevant to the political scientist. It requires the expertise of the political science scholars. Typically, actor dictionaries associate individuals with their groups and actions (Schrodt et al., 2008; Boschee et al., 2015). The high cost and human effort to update these repositories is unsustainable and requires continual updating. This is particularly problematic when an actor’s role changes or a new actor emerges. All of these complexities add to the costs and domain knowledge necessary for this component of extracting event data information from corpora.

3 Conflibert Details and Examples

Conflibert is a LLM trained on a *curated corpus of high-quality text data about politics, conflict, and violent events* from newswires and expert sources in English.⁶ Conflibert uses this domain knowledge with which it has been “trained” to be a more useful language model than vanilla BERT, LLM, or a simple dictionary approach. Conflibert is an engine or baseline for extracting information about political texts. It 1) sorts political violence texts from other ones (classification), 2) identifies possible political actors and entities (since it was trained to do so and does NER with this knowledge), and 3) provides for masking and question and answer (QA) tasks for coding. Clearly, the process and methodology here can be adapted to use other methods beyond BERT (other options and more recent LLMs are explored below). It can then be extended in domain areas / expertise, as well as scope conditions to include new languages, etc. via masking, fine-tuning, or other extensions.

Conflibert enables three distinct yet interconnected tasks that are crucial for a nuanced understanding of conflict dynamics: binary classification of violence-related news, multi-class classification of attack types, and named entity recognition in terrorism event reports. Users then take this model as given and a base against which improvements can be measured.⁷

⁶Training examples are listed at <https://github.com/eventdata/Conflibert/tree/main/pretrain-corpora> and test or evaluation ones are at <https://github.com/eventdata/Conflibert/tree/main/data>.

⁷The base Conflibert model is documented at <https://eventdata.utdallas.edu/> and <https://github.com/eventdata/>

An example of the first task is the binary classification of news articles to determine their relevance to gun violence. Utilizing a dataset comprising BBC news articles and the 20 Newsgroups corpus, we trained Conflibert to discern whether a given news item pertains to gun violence incidents. This fine-tuning task is significant for both domestic and international conflict studies, since it shows how to filter rapidly large volumes of news data about something like gun violence-related events. The ability to quickly identify relevant articles from a diverse news corpus can significantly enhance researchers' capacity to track and analyze (gun-related) conflicts in real-time.

For this binary classification task, Conflibert distinguishes between *gun violence-related* and *non-gun violence-related incidents*. Consider these examples:

Example

Input: "Two Lashkar e Jhangvi LeJ militants Asim alias Kapri and Ishaq alias Bobby confessed to killing four Rangers in Ittehad Town of Karachi, the provincial capital of Sindh."

Output: Gun Violence Related (1)

Input: "More than a week after a woman Communist Party of India-Maoist (CPI-Maoist) cadre was killed in an encounter in the forests of Lanjigarh block in Kalahandi District, the Maoists identified her as Sangita and called a bandh (general shutdown) in two Districts in protest against the killing."

Output: Gun Violence Related (1)

The second task expands on this binary classification to a more nuanced multi-class classification of attack types. Employing the Global Terrorism Database (GTD) to train Conflibert, it can classify attacks into nine distinct categories, including bombing/explosion, armed assault, assassination, and various forms of hostage-taking. Here are examples from the South Asia Terrorism Portal (SATP) dataset :

Conflibert-Manual. Those wishing to deploy the model can access it directly via <https://huggingface.co/eventdata-utd>. A user interface to evaluate short paragraphs and sentences is also available: <https://eventdata.utdallas.edu/conflibert-gui/> and <https://huggingface.co/spaces/eventdata-utd/Conflibert-Demo>. We then fine-tune the model simultaneously on all three tasks, using a weighted loss function that balances the importance of each task. To enhance the model's performance on individual tasks, we implement task-specific data augmentation techniques.

Example

Input: "Islamic State (IS) in the latest issue of its online magazine Dabiq claimed that the five of the nine Gulshan café attackers were suicide fighters... The mujahidin held a number of hostages as they engaged in a gun battle with apostate Bengali police and succeeded in killing and injuring dozens of disbelievers before attaining shahadah."

Output: Armed Assault

Input: "The ongoing construction work of an interstate bridge on Pranhita River on Maharashtra-Telangana border was thwarted by the Naxalites [Left Wing Extremists, LWEs] who set an excavator on fire and also damaged other equipment at the construction site at Gudem in Aheri taluka (revenue unit) of Gadchiroli District on April 26."

Output: Facility/Infrastructure Attack

Input: "Three boys sustained injuries when a landmine went off in Atmar Khel area of Baizai tehsil (revenue unit) in Mohmand Agency of Federally Administered Tribal Areas (FATA) on June 18."

Output: Bombing/Explosion

The third task ConflIBERT addresses is named entity recognition, crucial for extracting structured information from unstructured text, enabling more detailed and systematic analyses of conflict actors and targets. Using event reports (MUC-4), which contain annotations of terrorism events, we train ConflIBERT to identify and classify entities such as Perpetrator Organizations, Physical Targets, Victims, and Perpetrator Individuals. Here is an NER classification example using SATP:

Example

Input: "A senior Muttahida Qaumi Movement (MQM) [ORG] worker identified as Sohail Rasheed [PERSON], 30, was shot dead near his home in Naeemabad [LOC] in Korangi Town [LOC] of Karachi [LOC], the provincial capital of Sindh [LOC], on June 19 [DATE]."

Output:

Perpetrator Organization: Muttahida Qaumi Movement (MQM)

Victim: Sohail Rasheed Physical Target: Not specified

Location: Naeemabad, Korangi Town, Karachi, Sindh

Date: June 19

The versatility demonstrated across these tasks also suggests potential applications in related fields such as international relations, security studies, and public policy. By providing a tool that can simultaneously

categorize events, identify key actors and targets, and filter relevant information from large text corpora, Conflibert offers a powerful means of analyzing the complex landscape of modern conflicts.

Recall, that Conflibert was trained in 2021 on data that at this point is nearly 3 years old (Hu et al., 2022). So a question is, how well does it do with more contemporaneous events and data? Consider the following example that has been processed using the interface at <https://eventdata.utdallas.edu/conflibert-gui/> or <https://huggingface.co/spaces/eventdata-utd/Conflibert-Demo>:

Former President Donald Trump, the 2024 presumptive Republican presidential nominee, was escorted off the stage by Secret Service after gunshots were fired at his rally in Butler, Pennsylvania. Mr. Trump was injured from the incident, with blood appearing on the right side of his face. This occurred two days before the start of the Republican National Convention in Milwaukee. The Butler County, Pennsylvania, district attorney told the Associated Press that a shooter was dead and a rally attendee was killed.⁸

The outputs for each of the coding tasks are:

Binary Classification for Political Violence "Positive: The text is related to conflict, violence, or politics. (Confidence: 99.85%)"

NER entries returned are:

Location: "butler , pennsylvania", "milwaukee", "butler county , pennsylvania"

Organisation: "secret service", "republican national convention", "the associated press", "district"

Person: "former president donald trump , the 2024 presumptive republican presidential nominee", "attorney"

Multilabel Classification "Armed Assault (Confidence: 98.40%) Bombing or Explosion (Confidence: 5.39%) Kidnapping (Confidence: 0.44%) Other (Confidence: 0.95%)".

The only notable error is the inability to recognize "district attorney" as a person. Such an error can easily be corrected with additional fine tuning about legal actors and titles.

The implications of this multi-faceted approach to conflict analysis using Conflibert are far-reaching for political science research. By addressing multiple aspects of conflict analysis simultaneously, researchers can process diverse datasets more efficiently, from broad news corpora to specialized conflict databases. The scalability of this approach is particularly noteworthy. Once trained on these diverse tasks, Conflibert can be rapidly deployed to process large volumes of new data, enabling real-time or near-real-time analysis of emerging conflicts. This capability is invaluable for researchers and policymakers who need to quickly assess and respond to evolving conflict situations.

⁸"Former President Donald Trump Removed From Stage After Shots Fired at Pennsylvania Rally", CSPAN, July 13, 2024. Accessed 2024-09-09.

4 Evaluating Conflibert vs. other LLMs

Hu et al. (2022) established Conflibert’s superiority against a baseline of BERT (Devlin et al., 2018). They show that for binary classification and NER (the main tasks discussed here) Conflibert was better than BERT (based on using cased and uncased models) using F_1 and macro F_1 statistics for weighted precision and recall. In prior analyses, it scored better than BERT uncased and cased models, in multiple tests and across datasets of interest to political conflict scholars (general news reports, gun violence reports, terrorist attack accounts [by region], etc.)⁹ Hu et al. (2022, see , Table 3 and Figure 1) Häffner et al. (2023) and Croicu (2024) give additional and independent evidence of the model’s strong performance relative to known alternatives for different conflict texts and related tasks.

The focus here is on Conflibert’s efficacy in the two critical NLP tasks of binary classification (BC) and named entity recognition (NER) *compared to recent developments like generative AI LLMs*. Gauging Conflibert’s comprehension and extraction of information from conflict-related texts can be benchmarked against more recently created baselines from much large LLMs like Gemma 2, LLama 3.1, and Qwen 2.5. The goal is to assess the quality of an LLM like Conflibert and compare it to larger, more costly, and more computationally expensive alternatives.

We do this initially for two datasets that were used in the earlier comparisons of Conflibert to BERT: the BBC News Dataset and re3d.¹⁰ The BBC News dataset is used for the binary classification task (Greene and Cunningham, 2006) and consists of 2,225 news articles, with 1,490 records for training and 735 for testing. The articles cover five categories: business, entertainment, politics, sport, and tech. For the conflict classification task, the dataset articles are relabeled as either conflict-related (1) or not conflict-related (0) by expert coders who analyzed each article’s content and context. This dataset is particularly relevant providing a diverse range of news articles, thus testing Conflibert’s performance in sorting conflict-related content across various domains. The binary classification task mimics real-world scenarios where analysts must quickly identify political conflict-relevant information from a stream of news articles.

Once such articles or reports are identified via binary classification, then political actor and action classification are the relevant NER tasks. To compare Conflibert and more recent alternatives on this task, the re3d is used (Relationship and Entity Extraction Evaluation Dataset <https://github.com/dstl/re3d/>). These data are specifically designed for defense and security intelligence analysis, focused on the

⁹The diversity of comparisons include the BBC News dataset (Greene and Cunningham, 2006), a sample SATP dataset , the 20 Newsgroups dataset (Lang, 1995), the Gun Violence dataset (Pavlick et al., 2016), and the Event Status dataset (Huang et al., 2016) for binary classification tasks. For NER tasks, Conflibert CONT cased achieved the highest macro F1 score for the source and target labeling NER task on the CAMEO Codebook (Gerner et al., 2002) dataset, while Conflibert SCR uncased showed the highest macro F1 scores on both the MUC-4 (Defense Advanced Research Projects Agency Software and Intelligent Systems Technology Office, 1992) and the (DSTL, 2018) datasets.

¹⁰Both datasets are available for public use and can be accessed through the Conflibert GitHub repository.

conflict in Syria and Iraq, providing domain-specific content across various source and document types with differing entity densities. The entities of interest include organizations, persons, locations, and temporal expressions. Ground truth labels were established by annotators using a hybrid process.¹¹ The re3d is valuable to evaluate ConflBERT’s and LLMs’ extraction of relevant entities from conflict-related texts.

4.1 Methodology

Across these two datasets in this section, the performance of a task is done using versions of 1) ConflBERT, 2) Meta’s Llama 3.1 (8B), 3) Google’s Gemma 2 (9B), and 4) ConflLlama (8B). Note these are the most recent versions of these LLMs in late 2024. (We leave the comparisons to the much larger Qwen model to the next section). For the more recent LLMs, we offer brief descriptions:

Meta’s Llama 3.1 is the latest version of the Llama series of language models (Dubey et al., 2024). With 7 billion parameters, it strikes a balance between computational efficiency and performance.

Google’s Gemma 2 has 9 billion parameters, represents a significant advancement in the field of large language models (Team Gemma et al., 2024), offering robust performance across a wide range of NLP tasks while maintaining a relatively compact size.

Alibaba’s Qwen 2.5 has a large pre-training corpus focused on math and coding. Another key improvement, especially in the context that we are using the model for is the greater accuracy in generating structure outputs (as JSON objects).

ConflLlama based on LLaMA-3 8B, was specifically fine-tuned on the Global Terrorism Database (GTD) using QLoRA with a learning rate of 2e-4 and LoRA rank of 8. The model was trained with gradient checkpointing enabled and 4-bit quantization, achieving convergence with loss reduction from 1.95 to approximately 0.90 (Meher, 2024). We employ both $Q4_{K_M}$ and $Q8_0$ quantizations for comprehensive performance analysis.

Various performance metrics quantify how well the models classify an event or its key attributes (e.g., actors, actions, locations, and dates). These metrics essentially compare the ground truth with what the machine extracts to produce a numerical result. This *distance* between the two demonstrates the degree of congruence, and the goal for event data scientists is to achieve 100% congruence across multiple possible sources of error (Brandt and Sianan, 2024; Althaus et al., 2022). For binary classification the precision,

¹¹Described at <https://github.com/dstl/re3d/>. A modified version of the re3d, which underwent several preprocessing steps is used: These include tokenization and minor cleaning, removal of entity labels with confidence scores below 0.5, and resolution of overlapping entities by keeping only the largest span. The dataset was then converted to CoNLL 2003 format for compatibility with standard NER evaluation tools.

recall, and the F_1 score are reported. The focus here is the F_1 statistic, the geometric mean of the precision and recall of the classifications, combining both attributes. The NER tasks are evaluated for entity-level using precision, recall, macro F_1 , and span-level exact matches. The Macro F_1 used here measures this over the multiple attributes of the NER task.

4.2 Binary Classification

Table 1 shows the binary classification performance on the BBC News and re3d texts. ConflBERT has high precision for non-conflict texts (a low false positive rate) and high recall for conflict-related texts, suggesting a strong ability to identify relevant content. ConflBERT’s disparity between precision and recall for the conflict class indicates that it flags more texts as conflict-related. Gemma 2 shows perfect recall for non-conflict texts but fails to identify any conflict-related texts (zero precision, recall, and F1-scores for the conflict class). Meanwhile the Gemma 2 and Llama 3.1, while achieving a comparable overall accuracy, lack the nuanced understanding required for the specific task. Both have high recall and precision for the non-conflict texts but their performance on the conflict class remains poor and consequently have lower F1-scores. While Llama 3.1 is marginally better at detecting conflict-related content compared to Gemma 2, it still struggles significantly with this classification task.

Gemma 2 and Llama 3.1 show a bias towards classifying texts as non-conflict compared to ConflBERT — evident from their high precision and recall for the non-conflict class, coupled with poor performance on the conflict class. This imbalance suggests that general LLMs may overfit the majority class (non-conflict), potentially due to class imbalance in the training data or limitations in their ability to capture the nuanced features that distinguish conflict-related texts. While all three models achieve similar overall accuracy, ConflBERT demonstrates a more balanced and nuanced understanding of both classes, making it more suitable for real-world applications where identifying conflict-related content is crucial. The performance of Gemma 2 and Llama 3.1 shows that for a basic classification task, a domain-specific model that focuses on a local context is likely superior to a larger more general model when put to the same task. We turn to the issue of further fine-tuning the Llama, Gemma and related models below.

4.3 Named Entity Recognition Results

Table 2 shows the NER performance of the models for re3d. ConflBERT performs best across all key metrics with a weighted F1-score of 0.5981, compared to 0.3987 for Gemma 2 and 0.3809 for Llama 3.1. Noteworthy is ConflBERT’s performance in identifying non-entity tokens (class O) with an F1-score of 0.7493, higher than Gemma 2 (0.5709) and Llama 3.1 (0.5454). ConflBERT is more robust distinguishing between entity

Model	Class	Precision	Recall	F1-score	Support
Conflibert	0 (Non-conflict)	0.9827	0.8439	0.9080	269
	1 (Conflict)	0.5385	0.9245	0.6806	53
	Macro Avg	0.7606	0.8842	0.7943	322
	Weighted Avg	0.9096	0.8571	0.8706	322
Gemma 2 (9B)	0 (Non-conflict)	0.8354	1.0000	0.9103	269
	1 (Conflict)	0.0000	0.0000	0.0000	53
	Macro Avg	0.4177	0.5000	0.4552	322
	Weighted Avg	0.6979	0.8354	0.7605	322
Llama 3.1 (8B)	0 (Non-conflict)	0.8375	0.9963	0.9100	269
	1 (Conflict)	0.5000	0.0189	0.0364	53
	Macro Avg	0.6688	0.5076	0.4732	322
	Weighted Avg	0.7819	0.8354	0.7662	322

Table 1: Detailed Performance Metrics for Binary Classification

Model	Class	Precision	Recall	F1-score
Conflibert	Non-Entity (O)	0.7728	0.7271	0.7493
	Macro Avg	0.1871	0.1188	0.1407
	Weighted Avg	0.6510	0.5654	0.5981
Gemma 2 (9B)	Non-Entity (O)	0.7369	0.4660	0.5709
	Macro Avg	0.0030	0.0019	0.0023
	Weighted Avg	0.5146	0.3254	0.3987
Llama 3.1 (8B)	Non-Entity (O)	0.7335	0.4341	0.5454
	Macro Avg	0.0011	0.0006	0.0008
	Weighted Avg	0.5123	0.3032	0.3809

Table 2: Detailed Performance Metrics for Named Entity Recognition

and non-entity tokens, a fundamental task for political conflict text NER.

The precision and recall values for the non-entity class underscore Conflibert’s superiority. It maintains the balance between precision (0.7728) and recall (0.7271) compared to its counterparts. Both Gemma 2 and Llama 3.1 show a notable disparity between their precision and recall values, indicating potentially biased NER predictions, despite being larger models compared to Conflibert. This suggests that model size alone does not guarantee better performance in specialized NLP tasks like NER for political conflict texts. Conflibert’s architecture or training approach may be better suited for this specific task with more appropriate balance in precision and recall.

The macro and weighted average metrics provide insight into the models’ performance across all entity classes. Conflibert’s macro average F1-score (0.1407), while low, is higher than those of Gemma 2 (0.0023) and Llama 3.1 (0.0008). This indicates that Conflibert performs more consistently across different entity types, whereas the larger models struggle with certain classes. Conflibert demonstrates a clear advantage in this NER task, outperforming larger models across all metrics. Its balanced performance in identifying

Model	Task	Execution Time (s)	Max Memory (MB)	GPU Usage (%)
Conflibert	Classification	3.52	919.94	95.93
	NER	1.42	950.63	95.63
Llama 3.1 (8B)	Classification	575.23	950.64	92 - 94
	NER	489.39	950.65	92 - 94
Gemma 2 (9B)	Classification	730.14	950.65	90 - 97
	NER	866.23	957.19	90 - 97

Table 3: Performance Metrics for Conflibert, Llama 3.1, and Gemma 2 Models

non-entity tokens and its higher macro average scores suggest a more robust and versatile approach to named entity recognition. These results highlight the importance of specialized domain architectures and training strategies in tackling complex NLP tasks, rather than relying solely on model size.

4.4 Computational Performance Comparison

For both the BC task using the BBC News data and the NER task with the re3d, we performed experiments on a machine with GPU capabilities, and recorded the execution time, memory usage, CPU usage, and GPU usage for each LLM.¹² Table 3 presents the timings for each model and task combination. The most striking difference is the execution time: for classification, Conflibert took only 3.52 seconds, while Llama 3.1 (Gemma 2) took 575.23 (730.14) seconds. For NER, Conflibert completes the task in 1.42 seconds, compared to 489 (866) seconds for Llama 3.1 (Gemma 2). The speed of Conflibert can be attributed to its parallel architecture for processing input data. Conflibert’s superior performance stems from its ability to process inputs in parallel. BERT-based models, like Conflibert, can efficiently batch multiple inputs and process them simultaneously. Generative LLMs like Gemma and Llama, typically process inputs sequentially so each text requires a separate request to the model, introducing additional computational overhead. While we parallelize these models by batching multiple task requests, there are context-length constraints on processing the texts that differ across the models.

5 Classifying Texts about Terrorist Attacks

Some pre-training of the generative LLMs could bring their performance up to or exceeding the performance of Conflibert (see, Wang, 2024, commenting on Häffner et al. (2023)). Fine-tuning the LLMs like Conflibert permits an appropriate and critical comparison since learning additional parameters of neural layers

¹²These tests were conducted on a machine with an AMD Ryzen 7840HS Processor, NVIDIA GeForce RTX 4060 Laptop GPU 8GB GDDR6 (105W), and 32 GB 6400MHz LPDDR5X RAM. The software environment consisted of WSL (Ubuntu) with Python 3.10.12.

in the BERT representation of the texts in them is critical for additional extensions of models like ConflBERT. This asks for a comparison of ConflBERT to other more recent generative LLMs *with pre-training on political conflict texts*. Critical to this example is its replicability and service as a baseline comparison for event feature classification across new LLMs. This example replicates the common researcher problem: one has identified political conflict-related texts (or prior dataset to be extended) and organized them (say in a CSV, JSON, or other database) for analysis with standard NLP to extract the relevant event information. This then leaves open the choices of the LLM and the pre-training. As an illustration, consider the short texts in the Global Terrorism Dataset (GTD) (LaFree and Dugan, 2007).¹³ GTD is a good choice because it 1) is a comprehensive open-source database of terrorist events, 2) contains the conflict classification tasks (what kind of attack is in the event?), 3) provides consistent, well-structured texts for NLP tasks, and 4) is classified by experts: one knows from the codebook and the dataset who perpetrated the terrorist attacks, the nature of the attacks and the types of victims. One cannot use these texts for the BC task, but they are suitable for evaluating models' NER and event multi-label classifications.

The task here is predicting the categorization of terrorist attacks from each GTD text description, comparing the various LLMs' codings to the original (human) GTD annotations of the terrorist attack types. ConflBERT is compared to the aforementioned Llama and Gemma varieties, a larger LLM (Qwen 2.5), and a fine-tuned variant of Llama that we denote as ConflLlama.¹⁴ The training prompts for the generative LLMs are given in an appendix. The selection of evaluation metrics (ROC, accuracy, precision, recall, and F1-score) follows standard practices in conflict event classification (Schrodt and Van Brackle, 2020).

For testing and evaluation GTD data from 1970 to 2016 are used to train the LLMs and they are tested with data from 2017–2020. Most of the GTD events have no texts for 1970–1997, so this is mainly based on training texts from 1998–2016. The LLM coded texts produce sets of binary classifications of each of the nine GTD event types across 37,709 texts recorded in GTD using each of the six models (ConflBERT, ConflLlama4, ConflLlama8, Gemma, Llama, and Qwen). The first three of these are ones we produce, the latter three are "off the shelf" from `huggingface.co`.

5.1 Basic Classification Results

Figure 1 shows the comparative analysis of model performance differences across the LLM architectures. Here the left column presents receiver-operator characteristic curves (ROCs) for the conflict-trained LLMs while the right column presents the same for the general (non-conflict data trained), generative LLMs. The

¹³One consideration in selecting this example is the need for shareable and non-copyrighted texts for replication. This removes a barrier to entry for replication that would exist if data requiring a copyright, extensive downloads (scraping) or a license is used – e.g., Linguistic Data Consortium corpora. Further, we want something that is timely and relevant to conflict scholars. The short descriptions from GTD described here fit the bill.

¹⁴ConflLlama is a base Llama model given the training data and the classification prompt in the Appendix.

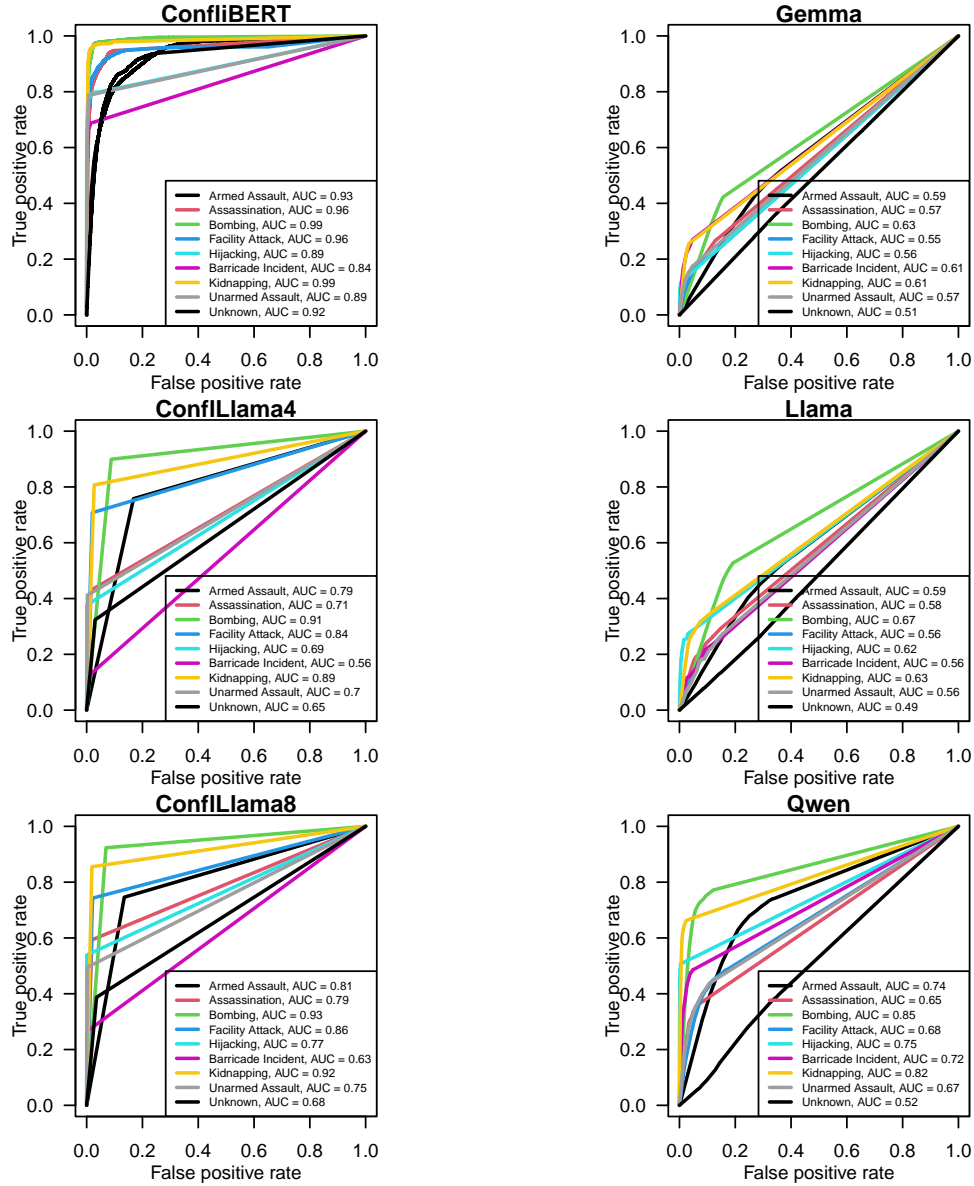


Figure 1: ROC and AUC for each LLM and event type. Curves along the northwestern edge are better.

results in the right column are closer to a 45° line, indicating nearly random classification performance by event-type. The area under the curve (AUC) for each event type are in the lower right.¹⁵ Across models, the higher accuracy of the ConfliBERT is evident and generally best for events about bombings and kidnappings (the green and gold lines) across the models – the most common kinds of attacks.

One criticism of only using an accuracy to compare models is that it is inflated by predicting the dominant class for imbalanced problems like the classifications here. Figure 2 shows the models’ precision-recall curves, in parallel with Figure 1. Best precision-recall curves are those that follow the top, northeastern edge of the plot.¹⁶ ConfliBERT has the highest precision-recall combinations for similar events (i.e., for the same colored

¹⁵AUC scores are relevant in conflict event classification due to balance precision and recall (Schrodt and Van Brackle, 2020)

¹⁶The numeric precision and recall scores commonly see in tables are weighted averages over the appropriate axes of these plots.

Model	Accuracy	Precision	Recall	F1	Total Time	Time / Document	Relative Speed
ConflBERT	0.90	0.83	0.77	0.79	27.6s	0.0016s	759.49x
ConflLlama-Q4KM* (8B)	0.72	0.72	0.72	0.71	49.9m	0.1746s	7.15x
ConflLlama-Q8* (8B)	0.76	0.76	0.76	0.75	52.3m	0.1831s	6.82x
Gemma 2 (9B)	0.60	0.27	0.19	0.21	3.1h	0.6605s	1.89x
Llama 3.1 (8B)	0.52	0.13	0.12	0.11	3.3h	0.7191s	1.74x
Qwen 2.5 (14B)	0.74	0.50	0.44	0.45	5.8h	1.2490s	1.00x

Table 4: Model Performance Comparison (Macro Averages). *ConflLlama timing measurements were performed on Delta HPC resources and are not directly comparable to other models’ timing metrics.

lines). Of the larger generative AI models, only the more recent and much larger Qwen model comes close to ConflBERT and ConflLlama in precision and recall performance, but only for kidnappings and bombings.

Precision-recall curves are a function of the cutoff used to classify a prediction as a match to GTD. Choosing the wrong cutoff, one may miss the benefits of a model to detect events (and mis-state their precision and recall in Figure 2). Figure 3 presents the F score for the precision-recall as a function of the chosen cutpoint for the correct classification *for each event type*. Unlike Figures 1 and 2, these are grouped by types of events, so the colors used indicate the models here. Ideally the values should be high across the cutpoints, like those for ConflBERT. ConflBERT and Qwen have the best F scores, followed by ConflLlama. These results align with previous findings suggesting that domain-specific fine-tuning often outperforms larger, general purpose models (Gururangan et al., 2020). Like in other specialized domains, ConflBERT’s strong performance can be attributed to its training on conflict-related data.

For the GTD conflict-related text analysis tasks, ConflBERT out performs the baseline competitors across all metrics shown in aggregate in Table 4. Its considerable speed improvements over larger models also reflect broader trends in NLP research emphasizing the importance of computational efficiency (Schwartz et al., 2020).¹⁷ For the fewer than 40K sentences evaluated here this is remarkably fast, yet for a larger document processing-training problem, the more general generative LLMs like Gemma, Llama, and Qwen are likely computationally prohibitive. While general-purpose large language models continue to improve, these results reinforce previous findings that specialized models can achieve superior performance in domain-specific tasks while maintaining significantly lower computational requirements (Strubell et al., 2019).

5.2 Multi-Label Classification Performance

Incidents that involve more than one event type are documented with multi-label classifications in the GTD. This occurs say when an incident includes an armed attack or assault in the course of a kidnapping. Multi-

¹⁷Processing times were measured on identical hardware configurations to ensure fair comparison.

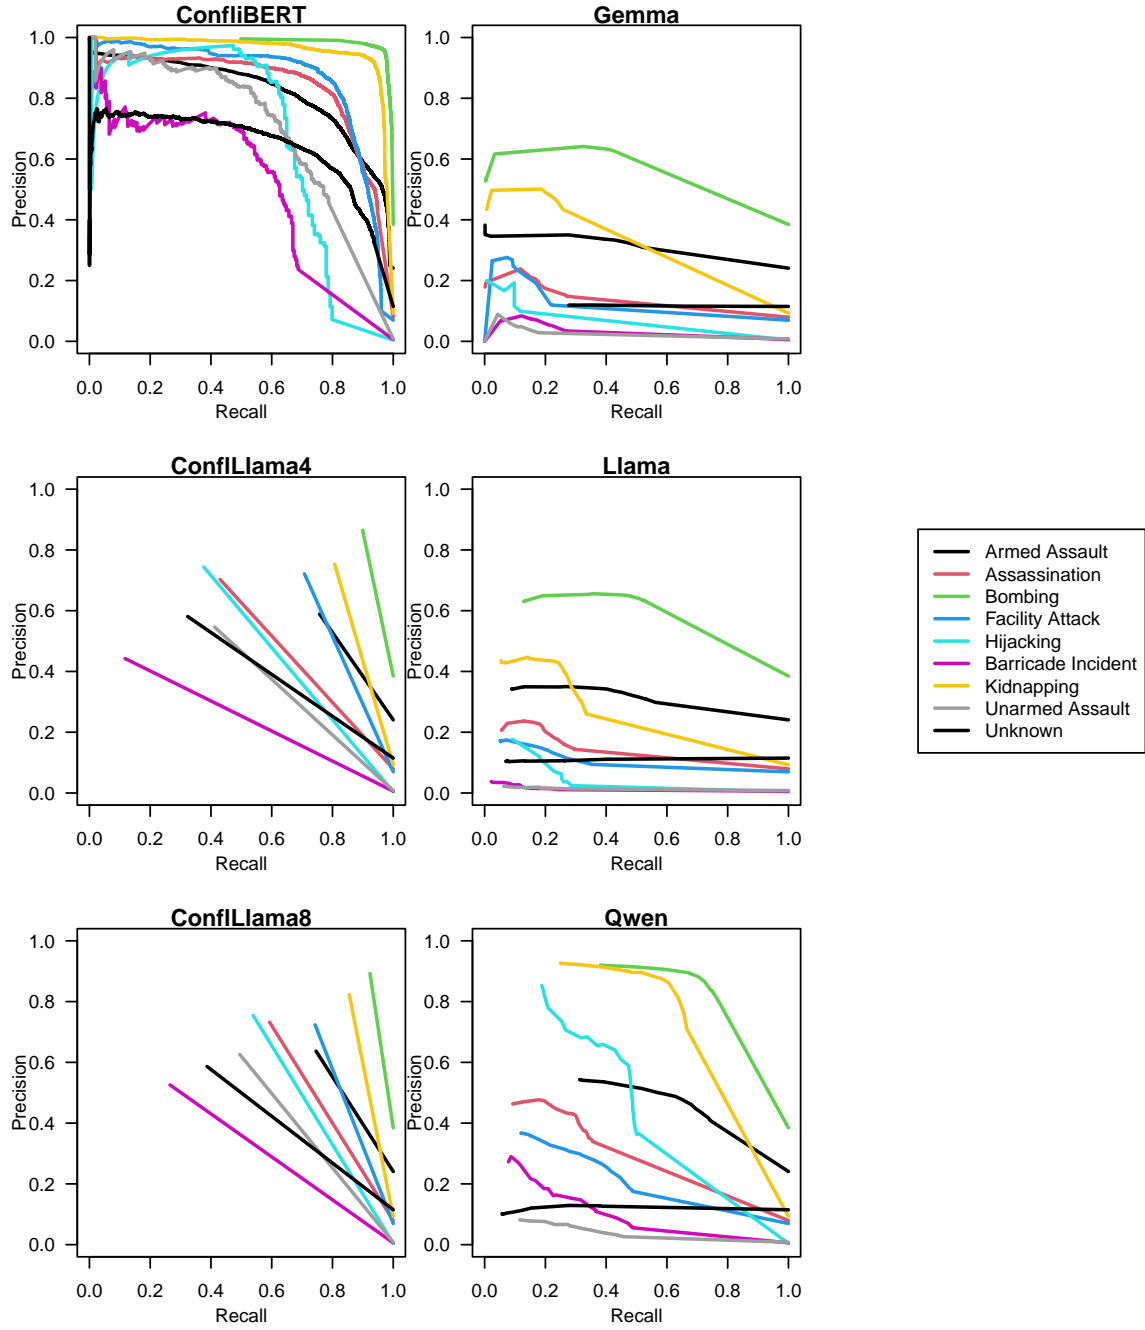


Figure 2: Precision-recall Curves for each LLM and event type. Curves along the northeastern edge are better.

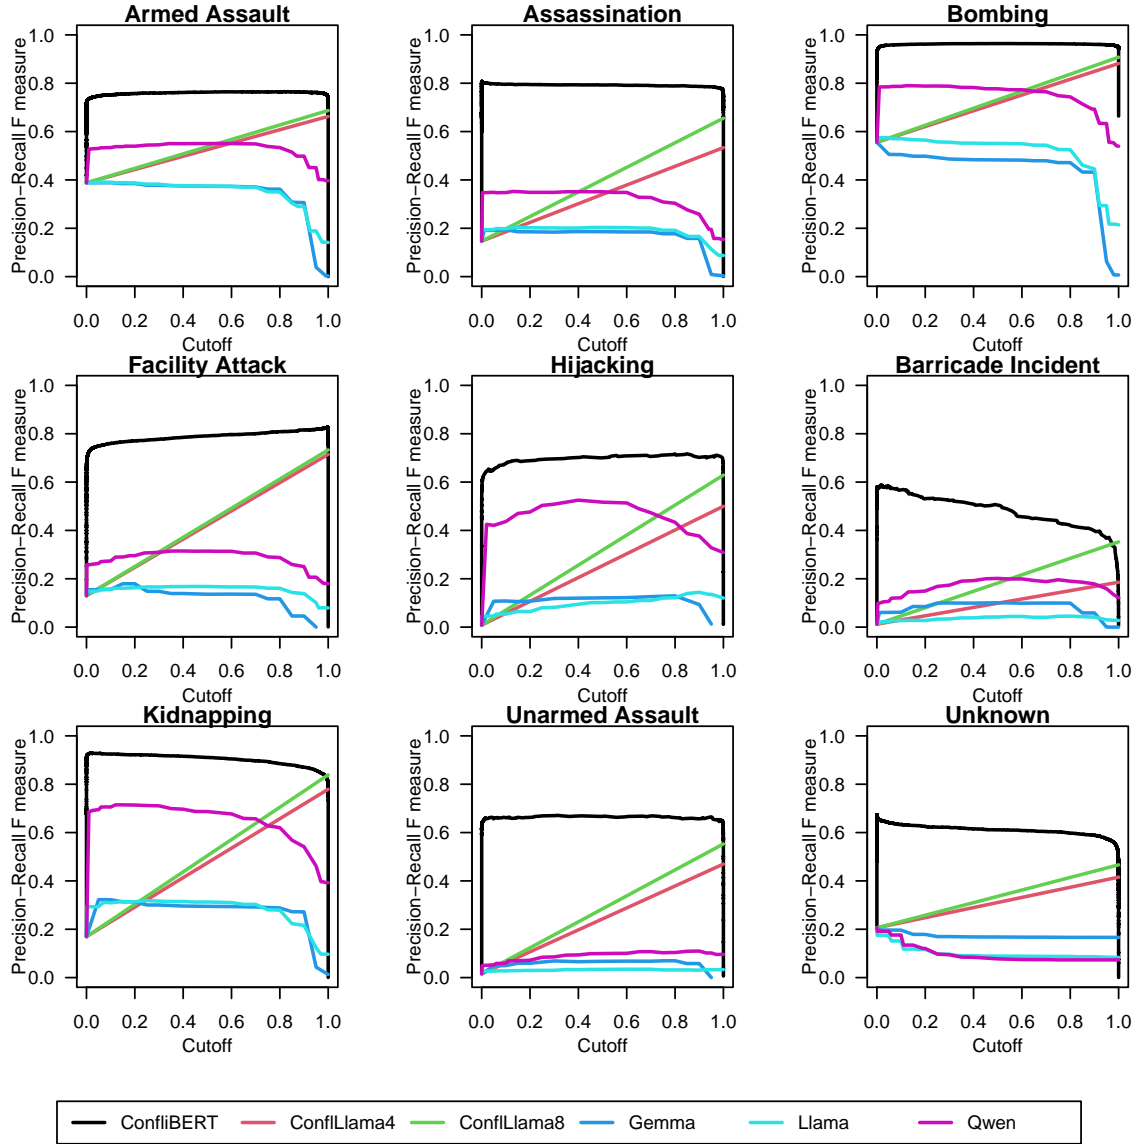


Figure 3: F-scores across cutoffs for each event type model. Higher curves are better.

Metric	ConfiBERT	ConfiLlama-Q8	ConfiLlama-Q4	Qwen	Gemma	LLaMA
Subset Accuracy (%)	79.38	72.40	68.80	50.99	30.70	32.03
Hamming Loss	0.035	0.052	0.061	0.096	0.133	0.148
Partial Match (%)	79.66	73.80	71.10	55.04	30.65	35.64
Label Cardinality						
True	0.963	0.963	0.963	0.963	0.963	0.963
Predicted	0.907	0.975	—	0.903	0.711	0.932

Table 5: Multi-Label Classification Metrics

label classification is important in conflict event coding, as real-world events often exhibit characteristics of multiple attack types (Radford, 2021). Less than 10% of the post-2016 (the test period) data has multi-label events. Originally, ConfiBERT was not built to handle this task, but can be used to accommodate it. Multi-label classification results, presented in Table 5, demonstrate ConfiBERT’s ability to handle complex event categorizations. The model achieved a subset accuracy of 79.38% and the lowest Hamming loss (0.035), indicating superior performance in scenarios where events may belong to multiple categories. The close alignment between predicted label cardinality (0.907) and true label cardinality (0.963) suggests that the model has effectively learned to capture the complexity of conflict events without over- or under-predicting multiple classifications.

The performance of ConfiBERT across all metrics suggests several important implications for conflict event classification. First, the results demonstrate that ConfiBERT with domain-specific fine-tuning can substantially outperform larger, general-purpose models, even when the latter have significantly more parameters (Gururangan et al., 2020). The model’s strong performance on rare event types is particularly noteworthy, as it addresses a common challenge in conflict event classification (Mueller and Rauh, 2021). This suggests that the fine-tuning process successfully captures the nuanced characteristics of different attack types, even with limited training examples.

5.3 Validity Comparisons

Another assessment of the classification differences from the LLMs is to consider how their distributions change over the event types. At any one point in (recent) time it may not be evident how the (mis-) classification of a given type of events affects inferences. But if there were systemic biases in LLM classification they are more evident as more types over events are collected — an inherently time series process for these data. This is particularly relevant in say a changepoint analysis of the drivers of trans-national terrorism like that addressed in Santifort et al. (2013, Figures 1–3), who use cumulative sums of terrorist event type classifications over time that would be severely biased upward or downward by LLM mis-classifications like

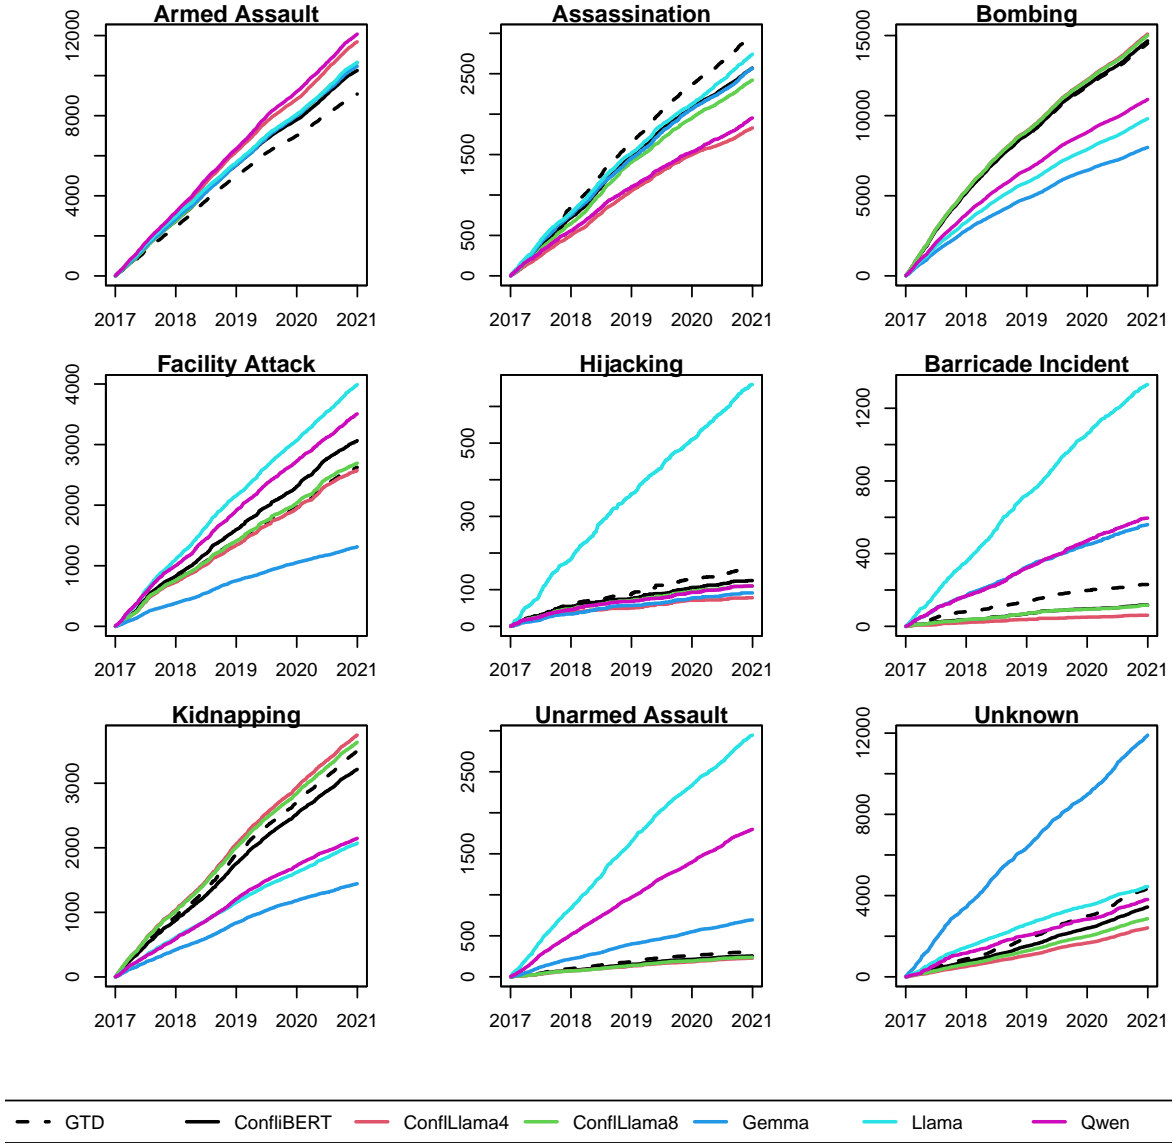


Figure 4: Cumulative number of predicted events, 2017–2021 by type and model

those documented above.

Figure 4 shows the cumulative time series of number of each type of GTD terrorist event from 2017–2020 as a dashed line. LLMs whose classifications are above this line are over-predicting / over-classifying the number of events of a given type while those under the dashed line are the reverse. A few immediate patterns jump out: the non-conflict pre-trained LLMs under classify bombing events (uppermost-right plot) – so Gemma, Llama, and Qwen. Second, the Llama, Qwen and Gemma models generally do poorly with the rarest event types (hijackings, barricade incidents, and unarmed assaults), but their performance includes over and under predictions relative to GTD’s human-coded data.

The reason such deviations and the relative performance of the LLMs matters here is that it will effect downstream time series analyses since systemtically mis-measured event types will lead to incorrect time

series dynamics and inferences that would confound those in works like Santifort et al. (2013). This builds on a key point since it shows that even using more data and more sophisticated methods for encoding texts, the issues of aggregation over time will still be important and affect inferences (Shellman, 2004).

6 Conclusions and Next Steps

Conflict research and event data have a fruitful history of incorporating NLP methods to advance methods of unstructured data processing (Schrodt (2001, 2012); Beiler et al. (2016)) and this work is a continuation of that path. The adoption of NLP techniques like those employed by ConflBERT improves how political scientists can extract and study events and political interactions. These tools offer the potential to analyze larger volumes of text data, enabling more comprehensive studies. They can reduce bias in event coding by applying consistent criteria across large datasets, identify patterns and trends that might be missed by human coders, and enable near-real-time analysis of political events as they unfold. While LLMs generally hold this potential, the approach in ConflBERT uses domain-specific knowledge, resulting in superior performance and even faster data processing of text classification and summarization tasks.

There are a series of conclusions to be drawn from this analysis. The results leverage existing infrastructure of BERT-alike LLMs and conflict researchers’ expertise to advance scholarship on conflict processes and international security. The contribution is that domain-specific knowledge – the things international and civil conflict scholars know – should be part of the information extraction process used to 1) filter relevant reports (binary classification) and 2) identify events, and 2) annotate their attributes (named entity recognition). A BERT-based model *plus* domain knowledge is able to do this in a way that is better on several metrics as documented in Section 4.

Beyond an infrastructure outline for political scientist to engage with texts about conflict and violence there are several other contributions of note here. First, ConflBERT builds on a known ontology (CAMEO/PLOVER) (Schrodt, 2012) for coding events and provides a set of tools for continuing to do so. This allows for additional fine-tuning of the models and a flatter development and learning curve. Unlike current large scale general LLMs, this allows researchers to openly and quickly work in this area (the span from ConflBERT in Hu et al. (2022) to the recent paper Osorio et al. (2024) is less than 36 months.)

Second, the typical social science conflict researcher need not build their own ConflBERT: one can fine-tune or extend this model since it is open and available for use via our website and HuggingFace <https://huggingface.co/eventdata-utd>. About 200GB of combined training data are invested in ConflBERT, ConflBERT Spanish, and ConflBERT Arabic. Additional classifications and training based on new ideas, texts, actors, etc. can be added and evaluated. We have done in this in the efforts to extend beyond event

coding just in English by working not just with a language and domain specific dictionary approach (Osorio and Reyes, 2017), but a general BERT-like model in Spanish (Yang et al., 2023). This shows how the domain-specific approaches can be brought to bear from say old codebooks and ontologies into the LLMs (Hu et al., 2024). So prior domain knowledge about regions, languages, and events can be part of how LLMs are used to encode and understand new texts and data.

Third, this approach is sometimes better than using larger LLMs. Unlike large generative LLMs, a decoder model like ConflBERT better fits what a social scientist needs, which is data extraction, organization, and (predictive) classification. Section 4 shows this in terms of performance metrics like accuracy, precision, F_1 , etc. It is also faster to use ConflBERT. While the initial LLM training for ConflBERT and its language variants took thousands of GPU hours, the work in Section 4.4 only takes hours of computing time on current laptops. Deploying this on a real data problem is scalable and feasible: it is 300-400 times faster than using a proprietary LLM for NER and 150-200 times faster for binary classification.

Finally, there are problems that can be addressed such as learning about and connecting events and actors. One area of interest is extending ontologies and NER to recognize and learn about new events and actors – who is the next leader, insurgent, or what are they doing? This is related to a literature on continual learning and catastrophic forgetting in LLMs. There is work in this area that can be applied and used to aid models like ConflBERT as well (e.g., Li et al., 2022). This would also be useful for extending text-as-data methods across networks of texts, languages, etc.

A Appendix: LLM prompts

ConflBERT and ConflLlama are fine-tuned specifically for terrorist event classification, without explicit prompting for output classifications given input event texts. For the general-purpose large language models (e.g., Gemma, Qwen, and Llama) the following structured prompt is used:

Example

Prompt: "Classify each of the following events into up to three of these categories, providing probabilities for each: Assassination, Armed Assault, Bombing/Explosion, Hijacking, Hostage Taking (Barricade Incident), Hostage Taking (Kidnapping), Facility/Infrastructure Attack, Unarmed Assault, Unknown

For each event, return only a JSON object with category names as keys and probabilities as values.

Example format: {"Armed Assault": 0.7, "Bombing/Explosion": 0.2, "Unknown": 0.1}

Events:"

This prompt follows key principles on effective LLM prompting Liu et al. (2023); Wei et al. (2022). Its structured format with explicit probability requirements builds on research showing that quantitative outputs improve model classification tasks (Brown, 2020). The multi-label approach, allowing up to three categories, reflects the complex classification task and the the original GTD structure – allowing direct comparisons. The JSON output format facilitates consistent parsing and evaluation, addressing challenges in systematic event coding. This standardization enables direct comparison with both human annotations and across models, while maintaining interpretability.

References

- Alsarra, S., Abdeljaber, L., Yang, W., Zawad, N., Khan, L., Brandt, P., Osorio, J., and D’Orazio, V. (2023). ConflBERT-Arabic: A pre-trained Arabic language model for politics, conflicts and violence. In *Proceedings of the 14th International Conference on Recent Advances in Natural Language Processing*, pages 98–108.
- Althaus, S., Peyton, B., and Shalmon, D. (2022). A Total Error Approach for Validating Event Data. *American Behavioral Scientist*, 66(5):603–624.
- Araci, D. (2019). FinBERT: Financial sentiment analysis with pre-trained language models.
- Beiler, J., Brandt, P. T., Halterman, A., Simpson, E., and Schrod, P. A. (2016). Generating Political Event Data in Near Real Time: Opportunities and Challenges. In Alvarez, R. M., editor, *Computational Social Science*. Cambridge University Press.
- Beltagy, I., Lo, K., and Cohan, A. (2019). SciBERT: A pretrained language model for scientific text. *arXiv preprint arXiv:1903.10676*.
- Boschee, E., Lautenschlager, J., O’Brien, S., Shellman, S., Starz, J., and Ward, M. (2015). ICEWS coded event data. *Harvard Dataverse*, 12.
- Brandt, P. T. and Sianan, M. (2024). Measurement in event data research. *Frontiers in Political Science – Political Science Methodologies*.
- Brown, T. B. (2020). Language models are few-shot learners. *arXiv preprint arXiv:2005.14165*.
- Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., and Androutsopoulos, I. (2020). LEGAL-BERT: the Muppets straight out of law school. *CoRR*, abs/2010.02559.
- Croicu, M. (2024). Deep active learning for data mining from conflict text corpora. *arxiv* <https://arxiv.org/abs/2402.01577>.
- Dai, Y., Radford, B., and Halterman, A. (2022). Political Event Coding as Text-to-Text Sequence Generation. In *Proceedings of the 5th Workshop on Challenges and Applications of Automated Extraction of Socio-political Events from Text (CASE)*, pages 117–123.
- Defense Advanced Research Projects Agency Software and Intelligent Systems Technology Office, editors (1992). *Fourth Message Understanding Conference (MUC-4): Proceedings of a Conference Held in McLean, Virginia, June 16-18, 1992*.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. (2018). BERT: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.
- DSTL (2018). Relationship and entity extraction evaluation dataset. Accessed: 2024-10-23.
- Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Yang, A., Fan, A., et al. (2024). The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.
- D’Orazio, V., Landis, S. T., Palmer, G., and Schrod, P. (2014). Separating the wheat from the chaff: Applications of automated document classification using support vector machines. *Political Analysis*, 22(2):224–242.
- Gerner, D., Jabr, R., and Schrod, P. (2002). Conflict and mediation event observations (CAMEO): A new event data framework for the analysis of foreign policy interactions. *International Studies Association, New Orleans*.
- Greene, D. and Cunningham, P. (2006). Practical solutions to the problem of diagonal dominance in kernel document clustering. In *Proceedings of the 23rd international conference on Machine learning*, pages 377–384.

- Grimmer, J., Roberts, M. E., and Stewart, B. M. (2022). *Text as data: A new framework for machine learning and the social sciences*. Princeton University Press.
- Gupta, T., Zaki, M., Krishnan, N. M. A., and Mausam (2022). MatSciBERT: A materials domain language model for text mining and information extraction. *npj computational materials*, 8(1):1–11.
- Gururangan, S., Marasović, A., Swayamdipta, S., Lo, K., Beltagy, I., Downey, D., and Smith, N. A. (2020). Don’t stop pretraining: Adapt language models to domains and tasks. *arXiv preprint arXiv:2004.10964*.
- Häffner, S., Hofer, M., Nagl, M., and Walterskirchen, J. (2023). Introducing an interpretable deep learning approach to domain-specific dictionary creation: A use case for conflict prediction. *Political Analysis*, 31(4):481–499.
- Halterman, A. (2017). Mordecai: Full Text Geoparsing and Event Geocoding. *Journal of Open Source Software*, 2(9):91.
- Halterman, A., Bagozzi, B. E., Beger, A., Schrodtt, P., and Scarborough, G. (2023a). PLOVER and POLE-CAT: A New Political Event Ontology and Dataset. *SocArXiv*.
- Halterman, A. and Radford, B. J. (2021). Few-Shot Upsampling for Protest Size Detection. *arXiv preprint arXiv:2105.11260*.
- Halterman, A., Schrodtt, P. A., Beger, A., Bagozzi, B. E., and Scarborough, G. I. (2023b). Creating custom event data without dictionaries: A bag-of-tricks.
- Hu, Y., Hosseini, M., Skorupa Parolin, E., Osorio, J., Khan, L., Brandt, P., and D’Orazio, V. (2022). ConflIBERT: A pre-trained language model for political conflict and violence. In Carpuat, M., de Marneffe, M.-C., and Meza Ruiz, I. V., editors, *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 5469–5482, Seattle, United States. Association for Computational Linguistics.
- Hu, Y., Skorupa Parolin, E., Khan, L., Osorio, J., and D’Orazio, V. (2024). Leveraging codebook knowledge with NLI and ChatGPT for zero-shot political relation classification. In Ku, L.-W., Martins, A., and Srikumar, V., editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 583–603, Bangkok, Thailand.
- Huang, K., Altosaar, J., and Ranganath, R. (2020). ClinicalBERT: Modeling clinical notes and predicting hospital readmission.
- Huang, R., Cases, I., Jurafsky, D., Condoravdi, C., and Riloff, E. (2016). Distinguishing past, on-going, and future events: The EventStatus corpus. In Su, J., Duh, K., and Carreras, X., editors, *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 44–54, Austin, Texas. Association for Computational Linguistics.
- Hui, B., Yang, J., Cui, Z., Yang, J., Liu, D., Zhang, L., Liu, T., Zhang, J., Yu, B., Lu, K., Dang, K., Fan, Y., Zhang, Y., Yang, A., Men, R., Huang, F., Zheng, B., Miao, Y., Quan, S., Feng, Y., Ren, X., Ren, X., Zhou, J., and Lin, J. (2024). Qwen2.5-coder technical report. *arxiv: https://arxiv.org/abs/2409.12186*.
- Imani, M. B., Chandra, S., Ma, S., Khan, L., and Thuraisingham, B. (2017). Focus Location Extraction from Political News Reports with Bias Correction. In *2017 IEEE International Conference on Big Data (Big Data)*, pages 1956–1964. IEEE.
- Imani, M. B., Khan, L., and Thuraisingham, B. (2019). Where Did the Political News Event Happen? Primary Focus Location Extraction in Different Languages. In *2019 IEEE 5th International Conference on Collaboration and Internet Computing (CIC)*, pages 61–70. IEEE.
- LaFree, G. and Dugan, L. (2007). Introducing the Global Terrorism Database. *Terrorism and Political Violence*.

- Lang, K. (1995). Newsweeder: Learning to filter netnews. In Prieditis, A. and Russell, S., editors, *Machine Learning Proceedings 1995*, pages 331–339. Morgan Kaufmann, San Francisco (CA).
- Lee, J.-S. and Hsiang, J. (2019). Patentbert: Patent classification with fine-tuning a pre-trained bert model.
- Lefebvre, C. and Stoeck, N. (2022). Rethinking the event coding pipeline with prompt entailment. *arXiv preprint arXiv:2210.05257*.
- Li, X., Wang, Z., Li, D., Khan, L., and Thuraisingham, B. (2022). LPC: A logits and parameter calibration framework for continual learning. In Goldberg, Y., Kozareva, Z., and Zhang, Y., editors, *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 7142–7155, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Liu, P., Zhang, L., and Gulla, J. A. (2023). Pre-train, prompt, and recommendation: A comprehensive survey of language modeling paradigm adaptations in recommender systems. *Transactions of the Association for Computational Linguistics*, 11:1553–1571.
- Meher, S. (2024). Confillama: Gtd-finetuned llama-3 8b. Based on Meta’s LLaMA-3 8B and GTD Dataset.
- Mueller, H. and Rauh, C. (2021). Forecasting civil conflict along the belt and road: A machine learning approach. *Journal of Peace Research*.
- Olsen, H., Simon, É., Velldal, E., and Øvrelid, L. (2024). Socio-political events of conflict and unrest: A survey of available datasets. In Hürriyetoglu, A., Tanev, H., Thapa, S., and Uludoğan, G., editors, *Proceedings of the 7th Workshop on Challenges and Applications of Automated Extraction of Socio-political Events from Text (CASE 2024)*, pages 40–53, St. Julians, Malta. Association for Computational Linguistics.
- Osorio, J., Alsarra, S., Converse, A., Alshammari, A., Heintze, D., Khan, L., Alatrush, N., Brandt, P. T., D’Orazio, V., Zawad, N., and Billah, M. (2024). Keep it local: Comparing domain-specific LLMs in native and machine translated text using parallel corpora on political conflict. In *2nd International Conference on Foundation and Large Language Models FLLM2024*, Dubai, UAE.
- Osorio, J. and Reyes, A. (2017). Supervised event coding from text written in Spanish: Introducing EVENTUS id. *Social Science Computer Review*, 35(3):406–416.
- Osorio, J., Reyes, A., Beltrán, A., and Ahmadzai, A. (2020). Supervised Event Coding from Text Written in Arabic: Introducing Hadath. In *Proceedings of the Workshop on Automated Extraction of Socio-political Events from News 2020*, pages 49–56.
- O’Connor, B., Stewart, B. M., and Smith, N. A. (2013). Learning to extract international relations from political context. *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics*, 1:1094–1104.
- Parolin, E. S., Hosseini, M., Hu, Y., Khan, L., Brandt, P. T., Osorio, J., and D’Orazio, V. (2022). Multi-CoPED: A Multilingual Multi-Task Approach for Coding Political Event Data on Conflict and Mediation Domain. In *Proceedings of the 2022 AAAI/ACM Conference on AI, Ethics, and Society*, pages 700–711.
- Parolin, E. S., Hu, Y., Khan, L., Osorio, J., Brandt, P. T., and D’Orazio, V. (2021). Come-ke: A new transformers based approach for knowledge extraction in conflict and mediation domain. In *2021 IEEE International Conference on Big Data (Big Data)*, pages 1449–1459. IEEE.
- Pavlick, E., Ji, H., Pan, X., and Callison-Burch, C. (2016). The gun violence database: A new task and data set for NLP. In Su, J., Duh, K., and Carreras, X., editors, *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 1018–1024, Austin, Texas. Association for Computational Linguistics.
- Radford, B. J. (2021). Automated dictionary generation for political eventcoding. *Political Science Research and Methods*, 9(1):157–171.

- Santifort, C., Sandler, T., and Brandt, P. T. (2013). Terrorist attack and target diversity: Changepoints and their drivers. *Journal of peace research*, 50(1):75–90.
- Schrodt, P. A. (2001). Automated coding of international event data using sparse parsing techniques. In *Annual Meeting of the International Studies Association, Chicago*. Citeseer.
- Schrodt, P. A. (2012). Precedents, progress, and prospects in political event data. *International Interactions*, 38(4):546–569.
- Schrodt, P. A. and Van Brackle, D. (2020). Automated coding of political event data. *Handbook of Research Methods in Political Science and International Relations*.
- Schrodt, P. A., Yilmaz, O., Gerner, D. J., and Hermreck, D. (2008). The CAMEO (Conflict and Mediation Event Observations) actor coding framework. In *2008 Annual Meeting of the International Studies Association*.
- Schwartz, R., Dodge, J., Smith, N. A., and Etzioni, O. (2020). Green AI. *Communications of the ACM*.
- Shellman, S. M. (2004). Time series intervals and statistical inference: The effects of temporal aggregation on event data analysis. *Political Analysis*, 12(1):97–104.
- Solaimani, M., Salam, S., Khan, L., Brandt, P. T., and D’Orazio, V. (2017a). APART: Automatic Political Actor Recommendation in real-Time. In Lee, D., Lin, Y.-R., Osgood, N., and Thomson, R., editors, *Social, Cultural, and Behavioral Modeling*, pages 342–348, Cham. Springer International Publishing.
- Solaimani, M., Salam, S., Khan, L., Brandt, P. T., and D’Orazio, V. (2017b). RePAIR: Recommend Political Actors in Real-time from News Websites. In *2017 IEEE International Conference on Big Data (Big Data)*, pages 1333–1340.
- Steinert-Threlkeld, Z. C. (2019). The future of event data is images. *Sociological Methodology*, 49(1):68–75.
- Strubell, E., Ganesh, A., and McCallum, A. (2019). Energy and policy considerations for deep learning in NLP. *arXiv preprint arXiv:1906.02243*.
- Team Gemma, Riviere, M., Pathak, S., Sessa, P. G., Hardin, C., Bhupatiraju, S., Hussenot, L., Mesnard, T., Shahriari, B., Ramé, A., et al. (2024). Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*.
- Wang, Y. (2024). On finetuning large language models. *Political Analysis*, 32(3):379–383.
- Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., et al. (2022). Emergent abilities of large language models. *arXiv preprint arXiv:2206.07682*.
- Wen, H., Lin, Y., Lai, T., Pan, X., Li, S., Lin, X., Zhou, B., Li, M., Wang, H., Zhang, H., Yu, X., Dong, A., Wang, Z., Fung, Y., Mishra, P., Lyu, Q., Surís, D., Chen, B., Brown, S. W., Palmer, M., Callison-Burch, C., Vondrick, C., Han, J., Roth, D., Chang, S.-F., and Ji, H. (2021). RESIN: A dockerized schema-guided cross-document cross-lingual cross-media information extraction and event tracking system. In Sil, A. and Lin, X. V., editors, *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies: Demonstrations*, pages 133–143, Online. Association for Computational Linguistics.
- Yang, W., Alsarra, S., Abdeljaber, L., Zawad, N., Delaram, Z., Osorio, J., Khan, L., Brandt, P. T., and D’Orazio, V. (2023). ConfiBERT-Spanish: A pre-trained Spanish language model for political conflict and violence. In *2023 7th IEEE Congress on Information Science and Technology (CiSt)*, pages 287–292. IEEE.